

CS229 Machine Learning Note 2

Classification and logistic regression

Terminologies

- **Classification Problems:** Predict y in small num of discrete values
- **Binary classification:** y takes 0, **negative class**, and 1, **positive class**, $y^{(i)}$ is also called **label** of the sample.

2.1 Logistic Regression

New Hypotheses

Since $y \in \{0, 1\}$, change form of hypotheses $h_\theta(x)$:

$$h_\theta(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

Still $\theta^T x = \theta_0 + \sum \theta_j x_j$

Logistic function or Sigmoid function:

$$g(z) = \frac{1}{1 + e^{-z}}$$

Note:

- $z \rightarrow \infty, g(z) \rightarrow 1, z \rightarrow -\infty, g(z) \rightarrow 0$ (This alternates discrete (0,1) into a continuous function on (0,1))
- Not hard to show:

$$g'(z) = \frac{1}{1 + e^{-z}} \left(1 - \frac{1}{1 + e^{-z}}\right) = g(z)(1 - g(z))$$

Estimate θ using MLE

Assume:

$$P(y|x; \theta) = h_\theta(x)^y (1 - h_\theta(x))^{1-y} \sim \text{Ber}(h_\theta(x))$$

$$\begin{aligned}\ell(\theta) &= p(\vec{y}|X; \theta) \\ &= \prod_{i=1}^n p(y^{(i)}|x^{(i)}; \theta) \\ &= \prod_{i=1}^n (h_\theta(x^{(i)}))^{y^{(i)}} (1 - h_\theta(x^{(i)}))^{1-y^{(i)}} \\ \ell(\theta) &= \sum_{i=1}^n y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)}))\end{aligned}$$

Optimize by $\theta := \theta + \alpha \nabla_\theta l(\theta)$: (Here is maximizing likelihood thus plus)

$$\begin{aligned}\frac{\partial}{\partial \theta_j} \ell(\theta) &= \left(y \frac{1}{g(\theta^T x)} - (1 - y) \frac{1}{1 - g(\theta^T x)} \right) \frac{\partial}{\partial \theta_j} g(\theta^T x) \\ &= \left(y \frac{1}{g(\theta^T x)} - (1 - y) \frac{1}{1 - g(\theta^T x)} \right) g(\theta^T x)(1 - g(\theta^T x)) \frac{\partial}{\partial \theta_j} \theta^T x \\ &= (y(1 - g(\theta^T x)) - (1 - y)g(\theta^T x))x_j \\ &= (y - h_\theta(x))x_j\end{aligned}$$

Therefore, **Stochastic Gradient Ascent Rule**:

$$\theta_j := \theta_j + \alpha(y^{(i)} - h_\theta(x^{(i)}))x_j^{(i)}$$

Note that

- it takes **Identical Form** to LMS.
- different from linear regression, LR has no normal function thus can only be optimized by gradient descent.

2.2 Digression: Perceptron(感知器) Learning Algorithm

Threshold function: modify Logistic Regression to discrete case.

$$g(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases} \quad h_{\theta}(x) = g(\theta^T x)$$

Keep update rule:

$$\theta_j := \theta_j + \alpha(y^{(i)} - h_{\theta}(x^{(i)}))x_j^{(i)}$$

- Perceptron regression hardly derives meaningful probabilistic interpretation
- So does deriving this by MLE

2.3 Multi-class classification

Multinomial distribution:

- $p(y|x; \theta)$ is a distribution over k possible discrete outcomes with $p = \phi_1, \dots, \phi_k$
- Introduce k groups of params $\theta_1, \dots, \theta_k$, then use $\theta_1^T x, \dots, \theta_k^T x$ to represent ϕ_1, \dots, ϕ_k
- Two problems
 - $\theta_j^T x$ is not necessarily 1.
 - Sum of $\theta_j^T x$ may not 1.

Softmax function: Addressing above, ensure output as probability vector.

$$\text{softmax}(t_1, t_2, \dots, t_k) = \begin{bmatrix} \frac{\exp(t_1)}{\sum_{j=1}^k \exp t_j} \\ \vdots \\ \frac{\exp(t_k)}{\sum_{j=1}^k \exp t_j} \end{bmatrix}$$

Logits: Inputs to the softmax function (vec t)

$$\Rightarrow \begin{bmatrix} P(y=1|x; \theta) \\ \vdots \\ P(y=k|x; \theta) \end{bmatrix} = \text{softmax}(\theta_1^T x, \dots, \theta_k^T x) = \begin{bmatrix} \frac{\exp(\theta_1^T x)}{\sum_{j=1}^k \exp \theta_j^T x} \\ \vdots \\ \frac{\exp(\theta_k^T x)}{\sum_{j=1}^k \exp \theta_j^T x} \end{bmatrix}$$

Written as

$$P(y = i|x; \theta) = \phi_i = \frac{\exp(t_i)}{\sum_{j=1}^k \exp(t_j)} = \frac{\exp(\theta_i^T x)}{\sum_{j=1}^k \exp(\theta_j^T x)}$$

$$\implies \ell(\theta) = \sum_{i=1}^n -\log \left(\frac{\exp(\theta_{y^{(i)}}^T x^{(i)})}{\sum_{j=1}^k \exp(\theta_j^T x^{(i)})} \right)$$

Cross-entropy Loss: $\ell_{ce} : \mathbb{R}^k \times \{1, \dots, k\} \rightarrow \mathbb{R}_{\geq 0}$ modularize in the complex equation above:

$$\ell_{ce}((t_1, \dots, t_k), y) = -\log \left(\frac{\exp(t_y)}{\sum_{j=1}^k \exp(t_j)} \right)$$

$$\ell(\theta) = \sum_{i=1}^n \ell_{ce}((\theta_1^T x^{(i)}, \dots, \theta_k^T x^{(i)}), y^{(i)})$$

By basic calculus, let $t = (t_1, \dots, t_k)$

$$\frac{\partial \ell_{ce}(t, y)}{\partial t_i} = \frac{\partial}{\partial t_i} (-t_y + \log(\sum_{j=1}^k \exp(t_j)))$$

$$= \phi_i - 1\{y = i\} \text{ or in vectorized notations } \phi - e_y$$

By chain rule:

$$\frac{\partial}{\partial \theta_i} \ell_{ce}((\theta_1^T x, \dots, \theta_k^T x), y) = (\phi_i - 1\{y = i\})x$$

Thus, **gradient of loss:**

$$\frac{\partial \ell(\theta)}{\partial \theta_i} = \sum_{j=1}^n (\phi_i^{(j)} - 1\{y^{(j)} = i\})x^{(j)}, \quad \phi_i^{(j)} = \frac{\exp(\theta_i^T x^{(j)})}{\sum_{s=1}^k \exp(\theta_s^T x^{(j)})}$$

2.4 Newton Method for Maximizing

Find $f(\theta) = 0$ by updating rule:

$$\theta := \theta - \frac{f(\theta)}{f'(\theta)}$$

$$\implies \theta := \theta - \frac{\ell'(\theta)}{\ell''(\theta)}$$

Newton-Raphson method Generalize to multidimensional setting:

$$\theta := \theta - H^{-1} \nabla_{\theta} \ell(\theta)$$

where $\nabla_{\theta} \ell(\theta)$ us vector of partial derivative of $\ell(\theta)$ to θ_i 's.

Hessian: d -by- d matrix, entry: $H_{ij} = \frac{\partial^2 \ell(\theta)}{\partial \theta_i \partial \theta_j}$ (黑塞矩阵)

Advantages

- Faster convergence (than gradient)
- Though more expensive per iteration. (When d is small, overall faster)
- **But perform worse at high dim problems** (as deriving inverse of H is required)

Applied to maximize logistic regression, also called **Fisher scoring**